

Q1. You have created an Elastic Load Balancing load balancer listening on port 80, and you registered it with a single Amazon Elastic Compute Cloud (Amazon EC2) instance also listening on port 80. A client makes a request to the load balancer with the correct protocol and port for the load balancer. In this scenario, how many connections does the balancer maintain?

Options

- Only single connection is maintained as both port numbers are same
- The load balancer maintains two separate connections: one connection with the client and one connection with the Amazon EC2 instance.
- Two connections for Amazon EC2 instance (one public and one private) and one connection with the client. Thus total 3
- Two connections for Amazon EC2 instance (one public and one private) and two connections with the client (one public and one private), hence 4.

Time Remaining: 02:58:02

[Submit Assessment](#)

Q2. You are a solutions architect working for a media company that hosts its website on AWS. Currently, there is a single Amazon Elastic Compute Cloud (Amazon EC2) Instance on AWS with MySQL installed locally to that Amazon EC2 Instance. You have been asked to make the company's production environment more resilient and to increase performance. You suggest that the company split out the MySQL database onto an Amazon RDS Instance with Multi-AZ enabled. This addresses the company's increased resiliency requirements. Now you need to suggest how you can increase performance. Ninety-nine percent of the company's end users are magazine subscribers who will be reading additional articles on the website, so only one percent of end users will need to write data to the site. What should you suggest to increase performance?

Options

- Alter the connection string so that if a user is going to write data, it is written to the secondary copy of the Multi-AZ database.
- Recommend that the company use read replicas, and distribute the traffic across multiple read replicas.
- Alter the connection string so that if a user is going to write data, it is written to the primary copy of the Multi-AZ database.
- Migrate the MySQL database to Amazon Redshift to take advantage of columnar storage and maximize performance.



Q3. You create an Auto Scaling group in a new region that is configured with a minimum size value of 10, a maximum size value of 100, and a desired capacity value of 50. However, you notice that 30 of the Amazon Elastic Compute Cloud (Amazon EC2) instances within the Auto Scaling group fail to launch. Which of the following is the cause of this behavior?

Options

- You cannot define an Auto Scaling group larger than 20.
- You did not attach an Elastic Load Balancing load balancer to the Auto Scaling group.
- The Auto Scaling group maximum value cannot be more than 20.
- You have not raised your default Amazon EC2 capacity (20) for the new region



Q4. 27. Currently, you're helping design and architect a highly available application. After building the initial environment, you discover that a part of your application does not work correctly until port 443 is added to the security group. After adding port 443 to the appropriate security group, how much time will it take before the changes are applied and the application begins working correctly? Choose the correct answer from the options below.

Options

- Generally, it takes 2-5 minutes in order for the rules to propagate.
- Changes apply instantly to the security group, and the application should be able to respond to 443 requests.
- Immediately after a reboot of the EC2 Instances belong to that security group.
- It will take 60 seconds for the rules to apply to all Availability Zones within the region.



Q5. An application in AWS is currently running in the Singapore region. You have been asked to implement disaster recovery for the same. So, if the application goes down in the Singapore region, it has to be started in the Asia region. Your application relies on pre-built AMIs. As a part of your disaster recovery strategy, which of the below points would you consider?

Options

- Nothing, because all AMIs by default are available in any region as long as they are created within the same account.
- Copy the AMI from the Singapore region to the Asia region. Modify the Auto Scalinggroups in the backup region to use the new AMI ID in the backup region.
- Modify the image permissions and share the AMI to the Asia region.
- Modify the image permissions to share the AMI with another account, then set the default region to the backup region.

Q6. Which of the following are the minimum required elements to create an Auto Scaling launch configuration?

Options

Launch configuration name, AMI, and instance type are needed to create an Auto Scaling launch configuration. Identifying a key pair, security group, and a block device mapping are optional elements for an Auto Scaling launch configuration.

Minimum a launch configuration name, AMI, security group and a block device mapping are needed to create an Auto Scaling launch configuration. Others are optional elements for an Auto Scaling launch configuration.

We need to have a minimum a launch configuration name, AMI, instance type and security group needed to create an Auto Scaling launch configuration. Identifying a key pair, and a block device mapping are optional elements for an Auto Scaling launch configuration.

Launch configuration name, AMI, instance type, key pair, security group, and block device mapping all are needed to create an Auto Scaling launch configuration.

