

## Database

## Legends



## Questions



01 02 03 04 05 06

Q1. 41. In Your company is hosting an application in AWS. The application is read intensive and consists of a set of web servers and AWS RDS. It has been noticed that the response time of the application increases due to the load on the AWS RDS instance. Which of the following measures can be taken to scale the data tier? Choose 2 answers from the options given below.

## Options

- Use Auto Scaling to scale out and scale in the database tier.
- Use SQS to cache the database queries.
- Use ElastiCache in front of your Amazon RDS DB to cache common queries.
- Use ElastiCache in front of your Amazon RDS DB to cache common queries.

[Mark for Review](#)[Prev](#)[Next](#)

Q2.

An

application requires a highly available relational database with an initial storage capacity of 8TB. This database will grow by 8GB everyday. <br>

To support the expected traffic, at least eight read replicas will be required to handle the database reads. Which of the below options meets these requirements? <br>

Options

DynamoDB

Amazon Aurora

Amazon S3

Amazon Redshift

Q3. An application consists of a web server and database server hosted on separate EC2 Instances. There are lot of read requests on the database which is degrading the performance of the application. Which of the following can help improve the performance of the database under this heavy load?

- | Options  |
|--|
| <input type="radio"/> Enable Multi-AZ for the database.                              |
| <input type="radio"/> Place another web server in the architecture to take the load. |
| <input checked="" type="radio"/> Put an ElastiCache in front of the database.        |
| <input type="radio"/> Place a CloudFront distribution in front of the database.      |

Q4. In Below are the requirements for a data store in AWS:

- a) Ability to perform SQL queries
- b) Integration with existing business intelligence tools
- c) High concurrency workload that generally involves reading and writing all columns of a small number of records at a time

Which of the following would be an ideal data store for the above requirements? Choose 2 answers from the options below.

A. AWS Redshift

B. AWS RDS

C. AWS Aurora

D. AWS S3

Options

B and C

C and D

A and B

A and C

Time Remaining: 02:50:24

Q5. A company is developing a web application to be hosted in AWS. This application needs a data store for session data.

As an AWS Solution Architect, which of the following would you recommend as an ideal option to store session data? Choose 2 answers from the options given below.

A. CloudWatch

B. DynamoDB

C. Elastic Load Balancing

D. ElastiCache

Options

A and B

B and C

B and D

C and D

Q6. Your company currently has data hosted in an Amazon Aurora MySQL DB. Since this data is critical, there is a need to ensure that it can be made available in another region in case of a disaster. How can this be achieved?

## Options

- Make a copy of the underlying EBS Volumes in the Amazon Cluster in another region.
- Creating a read replica of Amazon Aurora in another region.
- Enable Multi-AZ for the Aurora database.
- Create an EBS Snapshot of the underlying EBS Volumes in the Amazon Cluster and then copy them to another region.